

wrangle_act

December 28, 2019

1 WE NOW BEGIN OUR DATA WRANGLING JOURNEY.....

```
[1]: # We first import all the required libraries.
import pandas as pd
import numpy as np
import requests
import tweepy
from tweepy import OAuthHandler
import json
from timeit import default_timer as timer
```

1.1 GATHERING THE DATA

```
[2]: # We now import our provided twitter archive file with the help of pandas.
dog_ratings = pd.read_csv('twitter-archive-enhanced.csv')
```

```
[69]: pd.set_option("display.max_columns", 8)
dog_ratings.head()
```

```
[69]:
```

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
0	892420643555336193	NaN	NaN	
1	892177421306343426	NaN	NaN	
2	891815181378084864	NaN	NaN	
3	891689557279858688	NaN	NaN	
4	891327558926688256	NaN	NaN	

	timestamp	...	doggo	floofer	pupper	puppo
0	2017-08-01 16:23:56 +0000	...	None	None	None	None
1	2017-08-01 00:17:27 +0000	...	None	None	None	None
2	2017-07-31 00:18:03 +0000	...	None	None	None	None
3	2017-07-30 15:58:51 +0000	...	None	None	None	None
4	2017-07-29 16:00:24 +0000	...	None	None	None	None

[5 rows x 17 columns]

```
[4]: # Now we will programatically download the image prediction tsv file from
↳udacity's server using requests library.
```

```
url = 'https://d17h27t6h515a5.cloudfront.net/topher/2017/August/
↳599fd2ad_image-predictions/image-predictions.tsv'
response = requests.get(url)
with open("image_predictions.tsv", mode = 'wb') as outfile:
    outfile.write(response.content)
```

```
[71]: # We now read image predictions file into a df.
pd.set_option("display.max_columns", 8)
image_predictions = pd.read_csv('image_predictions.tsv', sep = '\t', encoding = '
↳utf-8')
image_predictions.head()
```

```
[71]:
```

	tweet_id	jpg_url	\
0	666020888022790149	https://pbs.twimg.com/media/CT4udn0WwAA0aMy.jpg	
1	666029285002620928	https://pbs.twimg.com/media/CT42GRgUYAA5iDo.jpg	
2	666033412701032449	https://pbs.twimg.com/media/CT4521TWwAEvMyu.jpg	
3	666044226329800704	https://pbs.twimg.com/media/CT5Dr8HUEAA-lEu.jpg	
4	666049248165822465	https://pbs.twimg.com/media/CT5IQmsXIAAKY4A.jpg	

	img_num	p1	...	p2_dog	p3	\
0	1	Welsh_springer_spaniel	...	True	Shetland_sheepdog	
1	1	redbone	...	True	Rhodesian_ridgeback	
2	1	German_shepherd	...	True	bloodhound	
3	1	Rhodesian_ridgeback	...	True	miniature_pinscher	
4	1	miniature_pinscher	...	True	Doberman	

	p3_conf	p3_dog
0	0.061428	True
1	0.072010	True
2	0.116197	True
3	0.222752	True
4	0.154629	True

[5 rows x 12 columns]

```
[6]: # Now we will gather retweet count and favourite count at minimum from twitter
↳API.
# Query Twitter API for each tweet in the Twitter archive and save JSON in a
↳text file
# These are hidden to comply with Twitter's API terms and conditions
consumer_key = 'P66aFQDJeBss9E0RraUg5CUCf'
consumer_secret = 'bUxnTS4JT5qfVJlfzdrSy8VHvFdjNhRLyLsbKlnn0qf3qdS7xm'
access_token = '760894336099688448-ECaKIkrsCGdyrHU56TEBzk3POPncpyk'
access_secret = 'QAnYI2sjEPsW8xyrYP4EUpbLHc5AGyGjMEuxP96nMyCM8'

auth = OAuthHandler(consumer_key, consumer_secret)
```

```
auth.set_access_token(access_token, access_secret)

api = tweepy.API(auth, wait_on_rate_limit=True)
```

[]:

```
[ ]: # Query Twitter's API for JSON data for each tweet ID in the Twitter archive
tweet_ids = dog_ratings.tweet_id.values
len(tweet_ids)

count = 0
fails_dict = {}
start = timer()
# Save each tweet's returned JSON as a new line in a .txt file
with open('tweet_json.txt', 'w') as outfile:
    for tweet_id in tweet_ids:
        count += 1
        print(str(count) + ": " + str(tweet_id))
        try:
            tweet = api.get_status(tweet_id, tweet_mode='extended')
            print("Success")
            json.dump(tweet._json, outfile)
            outfile.write('\n')
        except tweepy.TweepError as e:
            print("Fail")
            fails_dict[tweet_id] = e
            pass
end = timer()
print(end - start)
print(fails_dict)

# Hiding very large output of this cell.
%%hide output
```

[]:

```
[8]: # now we will open the file tweet_json.txt and read it line by line into the
      ↳ data frame.
# load json data into pandas DataFrame
text_file_path = 'tweet_json.txt'
tweet_json_df = pd.read_json(text_file_path, lines = True)
# extract columns that pertain to like and favorite counts
columns_of_interest = ['id', 'retweet_count', 'favorite_count']
tweet_likes = tweet_json_df[columns_of_interest]
```

1.2 ASSESSING THE DATA

```
[72]: # we will now visually assess dog_ratings df for Quality (Dirty Data) and
      ↪ tidiness(messy data) issues.
pd.set_option("display.max_columns", 8)
dog_ratings.head()
```

```
[72]:      tweet_id  in_reply_to_status_id  in_reply_to_user_id \
0  892420643555336193                NaN                NaN
1  892177421306343426                NaN                NaN
2  891815181378084864                NaN                NaN
3  891689557279858688                NaN                NaN
4  891327558926688256                NaN                NaN
```

```
      timestamp ... doggo floofer pupper puppo
0  2017-08-01 16:23:56 +0000 ...  None   None   None   None
1  2017-08-01 00:17:27 +0000 ...  None   None   None   None
2  2017-07-31 00:18:03 +0000 ...  None   None   None   None
3  2017-07-30 15:58:51 +0000 ...  None   None   None   None
4  2017-07-29 16:00:24 +0000 ...  None   None   None   None
```

[5 rows x 17 columns]

```
[73]: pd.set_option("display.max_columns", 8)
dog_ratings.sample(5)
```

```
[73]:      tweet_id  in_reply_to_status_id  in_reply_to_user_id \
301  836677758902222849                NaN                NaN
277  840370681858686976                NaN                NaN
184  856526610513747968      8.558181e+17      4.196984e+09
170  859074603037188101                NaN                NaN
902  758467244762497024                NaN                NaN
```

```
      timestamp ... doggo floofer pupper puppo
301  2017-02-28 20:41:37 +0000 ...  None   None   None   None
277  2017-03-11 01:15:58 +0000 ...  None   None   None   None
184  2017-04-24 15:13:52 +0000 ...  None   None   None   None
170  2017-05-01 15:58:40 +0000 ...  None   None   None   None
902  2016-07-28 01:00:57 +0000 ...  None   None   None   None
```

[5 rows x 17 columns]

```
[74]: # let us visually assess image_predictions for quality and tidiness issues.
pd.set_option("display.max_columns", 8)
image_predictions.head()
```

```
[74]:          tweet_id          jpg_url \
0  666020888022790149  https://pbs.twimg.com/media/CT4udn0WwAA0aMy.jpg
1  666029285002620928  https://pbs.twimg.com/media/CT42GRgUYAA5iDo.jpg
2  666033412701032449  https://pbs.twimg.com/media/CT4521TWwAEvMyu.jpg
3  666044226329800704  https://pbs.twimg.com/media/CT5Dr8HUEAA-lEu.jpg
4  666049248165822465  https://pbs.twimg.com/media/CT5IQmsXIAAKY4A.jpg
```

```
      img_num      p1 ... p2_dog      p3 \
0          1  Welsh_springer_spaniel ...   True   Shetland_sheepdog
1          1             redbone ...   True  Rhodesian_ridgeback
2          1      German_shepherd ...   True      bloodhound
3          1  Rhodesian_ridgeback ...   True  miniature_pinscher
4          1  miniature_pinscher ...   True      Doberman
```

```
      p3_conf  p3_dog
0  0.061428   True
1  0.072010   True
2  0.116197   True
3  0.222752   True
4  0.154629   True
```

[5 rows x 12 columns]

```
[75]: pd.set_option("display.max_columns", 8)
      image_predictions.sample(5)
```

```
[75]:          tweet_id          jpg_url \
210  669993076832759809  https://pbs.twimg.com/media/CUxLJ08U8AAu6Zu.jpg
1914  854120357044912130  https://pbs.twimg.com/media/C9px7jyVwAAnmwN.jpg
373   672988786805112832  https://pbs.twimg.com/media/CVbvjKqW4AA_CuD.jpg
371   672980819271634944  https://pbs.twimg.com/media/CVbodBOUsAAb7jZ.jpg
415   674014384960745472  https://pbs.twimg.com/media/CVqUgTIUAAUA8Jr.jpg
```

```
      img_num      p1 ... p2_dog      p3 \
210          1      piggy_bank ...  False      toy_poodle
1914         4  black-and-tan_coonhound ...   True      bluetick
373          1      Lakeland_terrier ...   True  wire-haired_fox_terrier
371          1      car_mirror ...   True      beagle
415          1      Pembroke ...   True      Eskimo_dog
```

```
      p3_conf  p3_dog
210  0.086502   True
1914  0.021762   True
373   0.038160   True
371   0.112397   True
415   0.068321   True
```

[5 rows x 12 columns]

```
[13]: # Let us visually assess tweet_likes for quality and tidiness issues.
tweet_likes.head()
```

```
[13]:
```

	id	retweet_count	favorite_count
0	892420643555336193	7840	36773
1	892177421306343426	5804	31666
2	891815181378084864	3843	23848
3	891689557279858688	8000	40095
4	891327558926688256	8650	38294

```
[14]: tweet_likes.sample(5)
```

```
[14]:
```

	id	retweet_count	favorite_count
2209	668221241640230912	192	502
1492	691090071332753408	339	1736
1367	700167517596164096	747	2666
1080	735274964362878976	5100	13156
1784	676916996760600576	1802	2979

```
[15]: # let us now also programmatically analyse data for quality issues.
dog_ratings.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null int64
in_reply_to_status_id   78 non-null float64
in_reply_to_user_id     78 non-null float64
timestamp               2356 non-null object
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id     181 non-null float64
retweeted_status_user_id 181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls           2297 non-null object
rating_numerator         2356 non-null int64
rating_denominator       2356 non-null int64
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                 2356 non-null object
puppo                   2356 non-null object
dtypes: float64(4), int64(3), object(10)
memory usage: 313.0+ KB
```

```
[16]: dog_ratings[dog_ratings.tweet_id.duplicated()]
```

```
[16]: Empty DataFrame
      Columns: [tweet_id, in_reply_to_status_id, in_reply_to_user_id, timestamp,
      source, text, retweeted_status_id, retweeted_status_user_id,
      retweeted_status_timestamp, expanded_urls, rating_numerator, rating_denominator,
      name, doggo, floofer, pupper, puppo]
      Index: []
```

```
[17]: dog_ratings.tweet_id.value_counts().sample(20)
```

```
[17]: 732732193018155009      1
      732585889486888962      1
      810284430598270976      1
      822163064745328640      1
      692417313023332352      1
      888554962724278272      1
      678255464182861824      1
      883117836046086144      1
      866334964761202691      1
      670003130994700288      1
      670668383499735048      1
      786595970293370880      1
      680940246314430465      1
      684188786104872960      1
      671486386088865792      1
      711363825979756544      1
      779124354206535695      1
      705442520700944385      1
      814638523311648768      1
      866686824827068416      1
      Name: tweet_id, dtype: int64
```

```
[18]: image_predictions.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id      2075 non-null int64
jpg_url       2075 non-null object
img_num       2075 non-null int64
p1            2075 non-null object
p1_conf       2075 non-null float64
p1_dog        2075 non-null bool
p2            2075 non-null object
p2_conf       2075 non-null float64
p2_dog        2075 non-null bool
```

```

p3          2075 non-null object
p3_conf     2075 non-null float64
p3_dog      2075 non-null bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB

```

```
[76]: pd.set_option("display.max_columns", 8)
      image_predictions.head()
```

```
[76]:
```

	tweet_id	jpg_url \
0	666020888022790149	https://pbs.twimg.com/media/CT4udn0WwAA0aMy.jpg
1	666029285002620928	https://pbs.twimg.com/media/CT42GRgUYAA5iDo.jpg
2	666033412701032449	https://pbs.twimg.com/media/CT4521TWwAEvMyu.jpg
3	666044226329800704	https://pbs.twimg.com/media/CT5Dr8HUEAA-lEu.jpg
4	666049248165822465	https://pbs.twimg.com/media/CT5IQmsXIAAKY4A.jpg

	img_num	p1	...	p2_dog	p3 \
0	1	Welsh_springer_spaniel	...	True	Shetland_sheepdog
1	1	redbone	...	True	Rhodesian_ridgeback
2	1	German_shepherd	...	True	bloodhound
3	1	Rhodesian_ridgeback	...	True	miniature_pinscher
4	1	miniature_pinscher	...	True	Doberman

	p3_conf	p3_dog
0	0.061428	True
1	0.072010	True
2	0.116197	True
3	0.222752	True
4	0.154629	True

```
[5 rows x 12 columns]
```

```
[20]: image_predictions[image_predictions.tweet_id.duplicated()]
```

```
[20]: Empty DataFrame
      Columns: [tweet_id, jpg_url, img_num, p1, p1_conf, p1_dog, p2, p2_conf, p2_dog,
      p3, p3_conf, p3_dog]
      Index: []
```

```
[21]: tweet_likes.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2333 entries, 0 to 2332
Data columns (total 3 columns):
id          2333 non-null int64
retweet_count  2333 non-null int64
favorite_count 2333 non-null int64
dtypes: int64(3)

```


memory usage: 54.8 KB

```
[22]: tweet_likes[tweet_likes.id.duplicated()]
```

```
[22]: Empty DataFrame
      Columns: [id, retweet_count, favorite_count]
      Index: []
```

2 QUALITY ISSUES

2.1 dog_rating table

1. column 'timestamp' includes '+0000' in the end which is of no use so we can strip that away.
2. 'retweeted_status_id', 'retweeted_status_user_id', 'retweeted_status_timestamp', are not required so they can be dropped.
3. Erroneous data types for following columns 'tweet_id', 'timestamp'
4. missing 'expanded_urls'
5. certain ids are having missing names in the 'name' column.
6. Remove the rows that belong to retweets.
7. Remove elements from name column of the table that do not represent name of a dog.

2.2 image_predictions table

1. dog breed names, certain have capital and certain have small letters, so it is better to make each of them lower case
2. erroneous datatype 'tweet_id', 'image_num'

2.3 tweet_likes table

1. erroneous datatype in column 'id'

3 TIDINESS ISSUES

1. the 'doggo', 'pupper', 'foofer' and 'puppo' columns can be just a single column named 'dog_stages'
2. The dog_ratings, image_predictions and tweet_likes represent the same observational unit so they they can all be merged as a single data frame.

```
[77]: pd.set_option("display.max_columns", 8)
      dog_ratings.sample()
```

```
[77]:      tweet_id  in_reply_to_status_id  in_reply_to_user_id  \
737  780858289093574656              NaN              NaN

      timestamp  ...  doggo  floofer  pupper  puppo
737  2016-09-27 19:54:58 +0000  ...  None    None    None    None

[1 rows x 17 columns]
```

```
[ ]:
```

4 CLEANING THE DATA

```
[24]: # In order to clean the data we first of all will create the copy of all the
      ↪ data frames in order to retain the unmodified dataset.
dog_ratings_clean = dog_ratings.copy()
image_pred_clean = image_predictions.copy()
tweet_likes_clean = tweet_likes.copy()
```

```
[78]: pd.set_option("display.max_columns", 8)
      dog_ratings_clean.sample(5)
```

```
[78]:
```

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
8961	674752233200820224	NaN	NaN	
2970	796759840936919040	NaN	NaN	
1724	680085611152338944	NaN	NaN	
337	832397543355072512	NaN	NaN	
399	824796380199809024	NaN	NaN	

	timestamp	...	rating_denominator	name	dog_stages	\
8961	2015-12-10 00:47:23 +0000	...	10	None	puppo	
2970	2016-11-10 17:02:03 +0000	...	10	Romeo	floofer	
1724	2015-12-24 18:00:19 +0000	...	10	by	doggo	
337	2017-02-17 01:13:34 +0000	...	10	Eevee	doggo	
399	2017-01-27 01:49:15 +0000	...	10	Bailey	doggo	

	stage_name
8961	None
2970	None
1724	None
337	None
399	None

```
[5 rows x 15 columns]
```

4.1 Note - First of all we will tackle the missing data and then tidiness and quality of data respectively.....

5 Missing Data

1. Missing 'expanded_urls' in dog_ratings table.
2. Certain ids have missing names in 'name' column in dog_ratings table.

Since we cannot obtain data for the above issues, we will leave them as it is

```
[ ]:
```

```
[26]: # Now let us test if all retweets are removed. For this there should be no non-
      ↪ null values in retweeted_status_id, retweeted_status_user_id,
      ↪ retweeted_status_timestamp.
      dog_ratings_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null int64
in_reply_to_status_id    78 non-null float64
in_reply_to_user_id      78 non-null float64
timestamp                2356 non-null object
source                  2356 non-null object
text                    2356 non-null object
retweeted_status_id      181 non-null float64
retweeted_status_user_id 181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls            2297 non-null object
rating_numerator         2356 non-null int64
rating_denominator       2356 non-null int64
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                  2356 non-null object
puppo                   2356 non-null object
dtypes: float64(4), int64(3), object(10)
memory usage: 313.0+ KB
```

Therefore our test is complete and we have removed all retweets.

6 Tidiness Issues

```
[ ]:
```

1. dog_ratings table : The 'doggo', 'pupper', 'floofer' and 'puppo' columns can be just a single column named 'dog_stages'

6.0.1 Define

The 'doggo', 'pupper', 'floofer' and 'puppo' columns can be just a single column named 'dog_stages'. For the purpose we will use pd.melt function of pandas

6.0.2 Code

```
[79]: pd.set_option("display.max_columns", 8)
      dog_ratings_clean.head(1)
```

```
[79]:      tweet_id  in_reply_to_status_id  in_reply_to_user_id  \
0  892420643555336193                NaN                NaN

      timestamp ... rating_denominator    name  dog_stages  \
0  2017-08-01 16:23:56 +0000 ...                10  Phineas    doggo

      stage_name
0          None

[1 rows x 15 columns]
```

```
[28]: dog_ratings_clean = pd.melt(dog_ratings_clean, id_vars = ['tweet_id',
↳ 'in_reply_to_status_id', 'in_reply_to_user_id', 'timestamp', 'source',
↳ 'text', 'retweeted_status_id', 'retweeted_status_user_id',
↳ 'retweeted_status_timestamp', 'expanded_urls', 'rating_numerator',
↳ 'rating_denominator', 'name'], var_name = 'dog_stages', value_name =
↳ 'stage_name')
```

```
[ ]:
```

6.0.3 TEST

```
[80]: pd.set_option("display.max_columns", 8)
dog_ratings_clean.sample(5)
```

```
[80]:      tweet_id  in_reply_to_status_id  in_reply_to_user_id  \
5178  817171292965273600                NaN                NaN
7676  797971864723324932                NaN                NaN
48    882992080364220416                NaN                NaN
5332  796125600683540480                NaN                NaN
7056  666071193221509120                NaN                NaN

      timestamp ... rating_denominator    name  dog_stages  \
5178  2017-01-06 00:49:53 +0000 ...                10  Tebow    pupper
7676  2016-11-14 01:18:12 +0000 ...                10   None    puppo
48    2017-07-06 15:58:11 +0000 ...                10  Rusty    doggo
5332  2016-11-08 23:01:49 +0000 ...                10   None    pupper
7056  2015-11-16 01:52:02 +0000 ...                10   None    pupper

      stage_name
5178          None
7676          None
48          None
5332          None
7056          None

[5 rows x 15 columns]
```

6.1 2. The dog_ratings , image_predictions and tweet_likes represent the same observational unit so they they can all be merged as a single data frame.

6.1.1 DEFINE

The dog_ratings , image_predictions and tweet_likes represent the same observational unit so they they can all be merged as a single data frame USING pandas pd.merge function.

6.1.2 CODE

```
[30]: master_df = pd.merge(dog_ratings_clean, image_pred_clean, on = 'tweet_id')
```

```
[81]: pd.set_option("display.max_columns", 8)
master_df.head()
```

```
[81]:
```

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
0	892420643555336193	NaN	NaN	
1	892420643555336193	NaN	NaN	
2	892420643555336193	NaN	NaN	
3	892420643555336193	NaN	NaN	
4	892177421306343426	NaN	NaN	

	timestamp	...	p2_dog	p3	p3_conf	p3_dog
0	2017-08-01 16:23:56 +0000	...	False	banana	0.076110	False
1	2017-08-01 16:23:56 +0000	...	False	banana	0.076110	False
2	2017-08-01 16:23:56 +0000	...	False	banana	0.076110	False
3	2017-08-01 16:23:56 +0000	...	False	banana	0.076110	False
4	2017-08-01 00:17:27 +0000	...	True	papillon	0.068957	True

[5 rows x 26 columns]

```
[32]: tweet_likes.head()
```

```
[32]:
```

	id	retweet_count	favorite_count
0	892420643555336193	7840	36773
1	892177421306343426	5804	31666
2	891815181378084864	3843	23848
3	891689557279858688	8000	40095
4	891327558926688256	8650	38294

```
[33]: # We will rename column name 'id' as 'tweet id' so that it can match and be
merged with other data frames.
tweet_likes_clean = tweet_likes_clean.rename(columns = {'id' : 'tweet_id'})
tweet_likes_clean.head()
```

```
[33]:
```

	tweet_id	retweet_count	favorite_count
0	892420643555336193	7840	36773
1	892177421306343426	5804	31666

2	891815181378084864	3843	23848
3	891689557279858688	8000	40095
4	891327558926688256	8650	38294

```
[82]: pd.set_option("display.max_columns", 8)
master_df_comp = pd.merge(master_df, tweet_likes_clean, on= 'tweet_id')
master_df_comp.head()
```

```
[82]:      tweet_id  in_reply_to_status_id  in_reply_to_user_id  \
0  892420643555336193                NaN                NaN
1  892420643555336193                NaN                NaN
2  892420643555336193                NaN                NaN
3  892420643555336193                NaN                NaN
4  892177421306343426                NaN                NaN
```

	timestamp	...	p3_conf	p3_dog	retweet_count	\
0	2017-08-01 16:23:56 +0000	...	0.076110	False	7840	
1	2017-08-01 16:23:56 +0000	...	0.076110	False	7840	
2	2017-08-01 16:23:56 +0000	...	0.076110	False	7840	
3	2017-08-01 16:23:56 +0000	...	0.076110	False	7840	
4	2017-08-01 00:17:27 +0000	...	0.068957	True	5804	

	favorite_count
0	36773
1	36773
2	36773
3	36773
4	31666

[5 rows x 28 columns]

7 Quality Issues

7.1 1. dog_rating table : Column ‘timestamp’ includes ‘+0000’ in the end which is of no use so we can strip that away.

7.1.1 Define

Column ‘timestamp’ includes ‘+0000’ in the end which is of no use so we can strip that away using pandas

7.1.2 Code

```
[35]: master_df_comp.timestamp = dog_ratings_clean.timestamp.str[:-5]
```

7.1.3 Test

```
[83]: pd.set_option("display.max_columns", 8)
      master_df_comp.head()
```

```
[83]:      tweet_id  in_reply_to_status_id  in_reply_to_user_id  \
0  892420643555336193                NaN                NaN
1  892420643555336193                NaN                NaN
2  892420643555336193                NaN                NaN
3  892420643555336193                NaN                NaN
4  892177421306343426                NaN                NaN

      timestamp ...  p3_conf  p3_dog  retweet_count  \
0  2017-08-01 16:23:56 +0000 ...  0.076110  False          7840
1  2017-08-01 16:23:56 +0000 ...  0.076110  False          7840
2  2017-08-01 16:23:56 +0000 ...  0.076110  False          7840
3  2017-08-01 16:23:56 +0000 ...  0.076110  False          7840
4  2017-08-01 00:17:27 +0000 ...  0.068957   True          5804

      favorite_count
0          36773
1          36773
2          36773
3          36773
4          31666

[5 rows x 28 columns]
```

```
[ ]:
```

7.2 2. dog_rating table : Remove the rows that belong to retweets.

7.2.1 Define

```
[84]: # So in order to remove the retweets, we have to remove the rows having any
      ↪ entries in the 'retweeted_status_id', 'retweeted_status_user_id',
      ↪ 'retweeted_status_timestamp' columns.
      # Following is the code for that.
      pd.set_option("display.max_columns", 8)
      master_df_comp.drop(master_df_comp[master_df_comp['retweeted_status_id'].
      ↪ notnull()== True].index, inplace = True)
      master_df_comp.drop(master_df_comp[master_df_comp['retweeted_status_user_id'].
      ↪ notnull()== True].index, inplace = True)
      master_df_comp.drop(master_df_comp[master_df_comp['retweeted_status_timestamp'].
      ↪ notnull()== True].index, inplace = True)
      master_df_comp.head()
```

```
[84]:      tweet_id  in_reply_to_status_id  in_reply_to_user_id  \
0  892420643555336193                NaN                NaN
1  892420643555336193                NaN                NaN
2  892420643555336193                NaN                NaN
3  892420643555336193                NaN                NaN
4  892177421306343426                NaN                NaN

      timestamp  ...  p3_conf  p3_dog  retweet_count  \
0  2017-08-01 16:23:56 +0000  ...  0.076110  False            7840
1  2017-08-01 16:23:56 +0000  ...  0.076110  False            7840
2  2017-08-01 16:23:56 +0000  ...  0.076110  False            7840
3  2017-08-01 16:23:56 +0000  ...  0.076110  False            7840
4  2017-08-01 00:17:27 +0000  ...  0.068957   True            5804

      favorite_count
0              36773
1              36773
2              36773
3              36773
4              31666

[5 rows x 28 columns]
```

7.3 3. Erroneous data types for following columns 'tweet_id', 'timestamp' and 'image_num'.

7.3.1 DEFINE

Erroneous data types for following columns 'tweet_id', 'timestamp' and 'image_num'.

7.3.2 CODE

```
[38]: master_df_comp.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 7948 entries, 0 to 8243
Data columns (total 28 columns):
tweet_id                7948 non-null int64
in_reply_to_status_id   92 non-null float64
in_reply_to_user_id     92 non-null float64
timestamp               7948 non-null object
source                 7948 non-null object
text                   7948 non-null object
retweeted_status_id      0 non-null float64
retweeted_status_user_id 0 non-null float64
retweeted_status_timestamp 0 non-null object
expanded_urls           7948 non-null object
rating_numerator         7948 non-null int64
```



```

rating_denominator      7948 non-null int64
name                    7948 non-null object
dog_stages              7948 non-null object
stage_name              7948 non-null object
jpg_url                 7948 non-null object
img_num                 7948 non-null int64
p1                      7948 non-null object
p1_conf                 7948 non-null float64
p1_dog                  7948 non-null bool
p2                      7948 non-null object
p2_conf                 7948 non-null float64
p2_dog                  7948 non-null bool
p3                      7948 non-null object
p3_conf                 7948 non-null float64
p3_dog                  7948 non-null bool
retweet_count           7948 non-null int64
favorite_count          7948 non-null int64
dtypes: bool(3), float64(7), int64(6), object(12)
memory usage: 1.6+ MB

```

Here we see that 'tweet_id' and 'img_num' can be of object datatype as they wouldnt require any calculation and 'timestamp' must be a datetime datatype.

```
[85]: master_df_comp['timestamp'] = pd.to_datetime(master_df_comp['timestamp'])
```

```
[40]: master_df_comp.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 7948 entries, 0 to 8243
Data columns (total 28 columns):
tweet_id                7948 non-null int64
in_reply_to_status_id   92 non-null float64
in_reply_to_user_id     92 non-null float64
timestamp                7948 non-null datetime64[ns]
source                  7948 non-null object
text                    7948 non-null object
retweeted_status_id     0 non-null float64
retweeted_status_user_id 0 non-null float64
retweeted_status_timestamp 0 non-null object
expanded_urls           7948 non-null object
rating_numerator         7948 non-null int64
rating_denominator       7948 non-null int64
name                    7948 non-null object
dog_stages              7948 non-null object
stage_name              7948 non-null object
jpg_url                 7948 non-null object
img_num                 7948 non-null int64
p1                      7948 non-null object

```

```

p1_conf          7948 non-null float64
p1_dog           7948 non-null bool
p2               7948 non-null object
p2_conf          7948 non-null float64
p2_dog           7948 non-null bool
p3               7948 non-null object
p3_conf          7948 non-null float64
p3_dog           7948 non-null bool
retweet_count     7948 non-null int64
favorite_count    7948 non-null int64
dtypes: bool(3), datetime64[ns](1), float64(7), int64(6), object(11)
memory usage: 1.6+ MB

```

```
[41]: master_df_comp['img_num'] = master_df_comp['img_num'].astype(object)
```

```
[42]: master_df_comp['tweet_id'] = master_df_comp['tweet_id'].astype(object)
```

```
[43]: master_df_comp.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 7948 entries, 0 to 8243
Data columns (total 28 columns):
tweet_id          7948 non-null object
in_reply_to_status_id  92 non-null float64
in_reply_to_user_id  92 non-null float64
timestamp         7948 non-null datetime64[ns]
source            7948 non-null object
text              7948 non-null object
retweeted_status_id  0 non-null float64
retweeted_status_user_id  0 non-null float64
retweeted_status_timestamp  0 non-null object
expanded_urls      7948 non-null object
rating_numerator    7948 non-null int64
rating_denominator  7948 non-null int64
name              7948 non-null object
dog_stages         7948 non-null object
stage_name         7948 non-null object
jpg_url            7948 non-null object
img_num            7948 non-null object
p1                7948 non-null object
p1_conf           7948 non-null float64
p1_dog            7948 non-null bool
p2                7948 non-null object
p2_conf           7948 non-null float64
p2_dog            7948 non-null bool
p3                7948 non-null object
p3_conf           7948 non-null float64
p3_dog            7948 non-null bool

```

```

retweet_count          7948 non-null int64
favorite_count         7948 non-null int64
dtypes: bool(3), datetime64[ns](1), float64(7), int64(4), object(13)
memory usage: 1.6+ MB

```

7.3.3 TEST

```
[44]: master_df_comp.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 7948 entries, 0 to 8243
Data columns (total 28 columns):
tweet_id              7948 non-null object
in_reply_to_status_id  92 non-null float64
in_reply_to_user_id   92 non-null float64
timestamp             7948 non-null datetime64[ns]
source                7948 non-null object
text                  7948 non-null object
retweeted_status_id    0 non-null float64
retweeted_status_user_id 0 non-null float64
retweeted_status_timestamp 0 non-null object
expanded_urls         7948 non-null object
rating_numerator       7948 non-null int64
rating_denominator     7948 non-null int64
name                  7948 non-null object
dog_stages            7948 non-null object
stage_name            7948 non-null object
jpg_url               7948 non-null object
img_num               7948 non-null object
p1                    7948 non-null object
p1_conf               7948 non-null float64
p1_dog                7948 non-null bool
p2                    7948 non-null object
p2_conf               7948 non-null float64
p2_dog                7948 non-null bool
p3                    7948 non-null object
p3_conf               7948 non-null float64
p3_dog                7948 non-null bool
retweet_count         7948 non-null int64
favorite_count         7948 non-null int64
dtypes: bool(3), datetime64[ns](1), float64(7), int64(4), object(13)
memory usage: 1.6+ MB

```

7.4 4. 'retweeted_status_id', 'retweeted_status_user_id', 'retweeted_status_timestamp', are not required so they can be dropped.

7.4.1 DEFINE

'retweeted_status_id', 'retweeted_status_user_id', 'retweeted_status_timestamp', are not required so they can be dropped since the retweets have now been removed.

7.4.2 CODE

```
[45]: master_df_comp.drop(columns= ['retweeted_status_user_id',  
    ↪ 'retweeted_status_id', 'retweeted_status_timestamp'], inplace = True)
```

7.5 5. dog breed names, certain have capital and certain have small letters, so it is better to make each of them lower case

7.5.1 DEFINE

dog breed names, certain have capital and certain have small letters, so it is better to make each of them lower case in 'p1', 'p2', 'p3' columns.

7.5.2 CODE

```
[86]: pd.set_option("display.max_columns", 8)  
master_df_comp.head(1)
```

```
[86]:      tweet_id  in_reply_to_status_id  in_reply_to_user_id  \  
0  892420643555336193                NaN                NaN  
  
      timestamp  ...  p3_conf  p3_dog  retweet_count  \  
0  2017-08-01 16:23:56+00:00  ...  0.07611  False          7840  
  
      favorite_count  
0              36773  
  
[1 rows x 28 columns]
```

```
[47]: master_df_comp['p1'] = master_df_comp['p1'].str.lower()
```

```
[48]: master_df_comp['p2'] = master_df_comp['p2'].str.lower()
```

```
[49]: master_df_comp['p3'] = master_df_comp['p3'].str.lower()
```

7.5.3 TEST

```
[87]: pd.set_option("display.max_columns", 8)  
master_df_comp.sample(5)
```

```
[87]:
```

	tweet_id	in_reply_to_status_id	in_reply_to_user_id	\
6837	672523490734551040	NaN	NaN	
3622	733460102733135873	NaN	NaN	
7058	671362598324076544	NaN	NaN	
3743	727524757080539137	NaN	NaN	
4408	705591895322394625	NaN	NaN	

	timestamp	...	p3_conf	p3_dog	retweet_count	\
6837	2015-12-03 21:11:09+00:00	...	0.061596	True	165	
3622	2016-05-20 00:51:30+00:00	...	0.017379	False	1298	
7058	2015-11-30 16:18:11+00:00	...	0.077301	False	291	
3743	2016-05-03 15:46:33+00:00	...	0.003941	True	1237	
4408	2016-03-04 03:13:11+00:00	...	0.035638	True	1164	

	favorite_count
6837	621
3622	4251
7058	1083
3743	4518
4408	3215


```
[5 rows x 28 columns]
```

```
[ ]:
```

7.6 6.Remove elements from name column of the table that do not represent name of a dog.

7.6.1 DEFINE

Remove elements from name column of the table that do not represent name of a dog

7.6.2 CODE

```
[51]: master_df_comp['name'] = master_df_comp['name'].replace( to_replace = ['a', 'u',
↳ 'an' , 'the', 'by', 'very'], value = 'None')
```

```
[52]: master_df_comp['name'].sample(25)
```

```
[52]: 2324      None
3477    Nollie
6824      Taz
6808    Norman
3330    Lenox
2189     Rory
8194     None
2202     Dale
1425    Buddy
```

```

378      Zoey
6326     None
6198     None
605      None
1536     Sky
6399    Tedders
3683     None
6428     None
4763    Cassie
560      None
318      None
1562     None
43       Koda
7989    Churlie
3695     None
4876    Brian
Name: name, dtype: object

```

7.6.3 TEST

```
[53]: master_df_comp['name']
```

```

[53]: 0      Phineas
      1      Phineas
      2      Phineas
      3      Phineas
      4      Tilly
      5      Tilly
      6      Tilly
      7      Tilly
      8      Archie
      9      Archie
     10      Archie
     11      Archie
     12      Darla
     13      Darla
     14      Darla
     15      Darla
     16    Franklin
     17    Franklin
     18    Franklin
     19    Franklin
     20      None
     21      None
     22      None
     23      None
     24      Jax

```

```
25      Jax
26      Jax
27      Jax
28     None
29     None
```

...

```
8214    None
8215    None
8216    None
8217    None
8218    None
8219    None
8220    None
8221    None
8222    None
8223    None
8224    None
8225    None
8226    None
8227    None
8228    None
8229    None
8230    None
8231    None
8232    None
8233    None
8234    None
8235    None
8236    None
8237    None
8238    None
8239    None
8240    None
8241    None
8242    None
8243    None
```

```
Name: name, Length: 7948, dtype: object
```

```
[54]: # Let us now store our cleaned dataframe in a csv file
      master_df_comp.to_csv('twitter_archive_master.csv')
```

7.7 OUR DATASET IS NOW CLEANED. LET US NOW TURN TOWARDS VISUALISATION AND INSIGHTS

8 VISUALISATION

```
[88]: pd.set_option("display.max_columns", 8)
      master_df_comp.head()
```

```
[88]:      tweet_id  in_reply_to_status_id  in_reply_to_user_id  \
0   892420643555336193                NaN                NaN
1   892420643555336193                NaN                NaN
2   892420643555336193                NaN                NaN
3   892420643555336193                NaN                NaN
4   892177421306343426                NaN                NaN
```

```
      timestamp  ...  p3_conf  p3_dog  retweet_count  \
0  2017-08-01 16:23:56+00:00  ...  0.076110  False            7840
1  2017-08-01 16:23:56+00:00  ...  0.076110  False            7840
2  2017-08-01 16:23:56+00:00  ...  0.076110  False            7840
3  2017-08-01 16:23:56+00:00  ...  0.076110  False            7840
4  2017-08-01 00:17:27+00:00  ...  0.068957   True            5804
```

```
      favorite_count
0              36773
1              36773
2              36773
3              36773
4              31666
```

```
[5 rows x 28 columns]
```

```
[56]: master_df_comp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7948 entries, 0 to 8243
Data columns (total 25 columns):
tweet_id          7948 non-null object
in_reply_to_status_id  92 non-null float64
in_reply_to_user_id  92 non-null float64
timestamp         7948 non-null datetime64[ns]
source            7948 non-null object
text              7948 non-null object
expanded_urls      7948 non-null object
rating_numerator    7948 non-null int64
rating_denominator  7948 non-null int64
name              7948 non-null object
dog_stages         7948 non-null object
stage_name         7948 non-null object
```



```

jpg_url          7948 non-null object
img_num          7948 non-null object
p1               7948 non-null object
p1_conf          7948 non-null float64
p1_dog           7948 non-null bool
p2               7948 non-null object
p2_conf          7948 non-null float64
p2_dog           7948 non-null bool
p3               7948 non-null object
p3_conf          7948 non-null float64
p3_dog           7948 non-null bool
retweet_count    7948 non-null int64
favorite_count   7948 non-null int64
dtypes: bool(3), datetime64[ns](1), float64(5), int64(4), object(12)
memory usage: 1.4+ MB

```

```

[57]: stage_max = master_df_comp['stage_name'].value_counts()
      stage_max_df = pd.DataFrame(stage_max)
      stage_max_df

```

```

[57]:      stage_name
      None      7632
      pupper    212
      doggo      73
      puppo     23
      floofer     8

```

```

[58]: breed_max = master_df_comp['p1'].value_counts().head()
      breed_max_df = pd.DataFrame(breed_max)
      breed_max_df

```

```

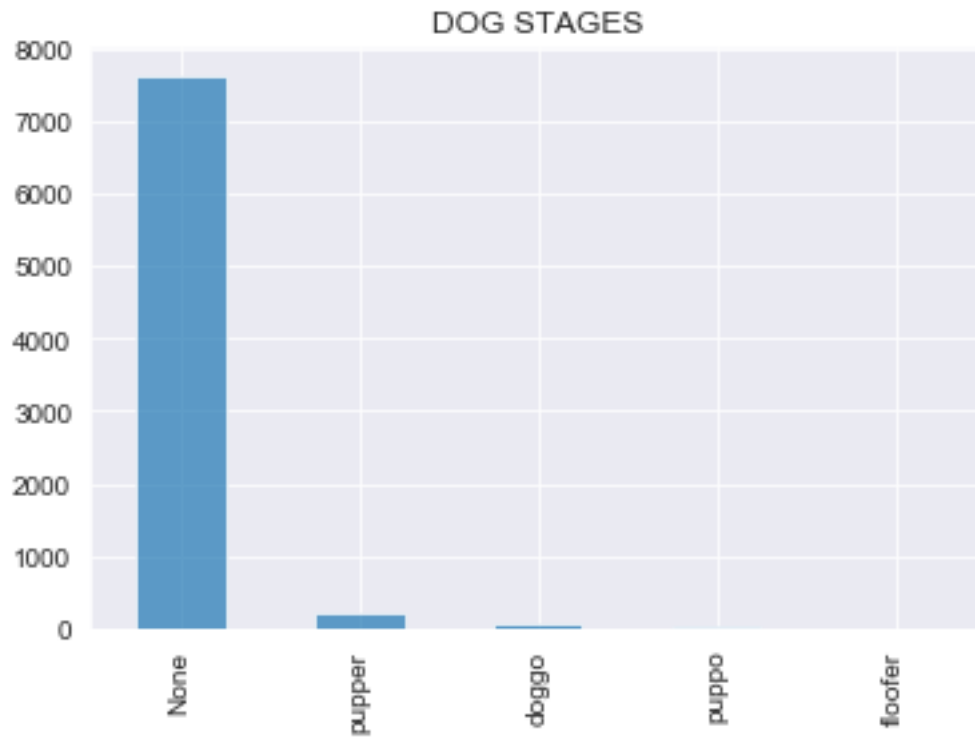
[58]:      p1
      golden_retriever    556
      labrador_retriever  372
      pembroke           352
      chihuahua           316
      pug                216

```


```

[62]: %matplotlib inline
      import seaborn as sns
      sns.set_style('darkgrid')
      stage_max_df.plot.bar( y = 'stage_name', alpha = 0.7, title = 'DOG STAGES',
      ↪ legend = '');

```



```
[63]: breed_max_df.plot.bar(y = 'p1', alpha= 0.7, legend = '', title = 'DOG BREED');
```



8.1 INSIGHTS

1. As we see from the above **visualisation no.1** that data wrangling specially for dog stages made it so convenient to see that most of the dogs could not be identified with a stage while maximum of the dogs whose stages could be defined lied in the pupper stage.
2. From the **second visualisation** we get that by **combining image_prediction table data together** made it easy to identify first prediction(p1) of the dog breed that came out to be golden retriever followed by labrador retriever and pembroke.
3. We see that by **removing the retweets** we are able to get a true picture of the data analysed as with retweets the counts would have been biased. we have also removed certain unrequired columns['retweeted_status_id', 'retweeted_status_user_id', 'retweeted_status_timestamp'] that contained retweet information.
4. We have also **altered the timestamp** to make it more presentable and also **changed the datatype of certain columns** ['tweet_id', 'timestamp' and 'image_num'] which would make the work of data analysts who further utilise the dataset easier.

I believe that above wrangling and insights are in accordance with the project motivation and have formed a clean dataset.